

Master-Practical Course: Scientific Computing
Computational Fluid Dynamics

Worksheet 4
Parallelization

Deadline: Tuesday 25th June, 2019, 12:00

Part I.

Introduction to Parallelization

1. Parallelization

In this worksheet, we address the parallelization of our flow simulation program. With parallelization we intend to reduce the total computing time by distributing the computational work among several processors, which perform the calculations concurrently. In addition, we often profit from a substantially larger amount of main memory compared to the memory that is available on sequential or stand-alone machines. This allows us to solve larger problems.

In this worksheet you will implement a distributed memory parallelization for the Navier-Stokes Equations

For the sake of simplicity, we restrict ourself to the parallelization of the algorithm presented in Worksheet 1.

2. Parallel Computers and Programming Environments

Before starting the exercises, we give a little overview on parallel computers and parallel programming. If you're familiar with all this, you're welcome to skip this section.

We start with a classification of possible (parallel) architectures. This very common classification was given by Flynn: It distinguishes between four basic types – SISD, MISD, SIMD, and MIMD – which can be explained as follows:

- Very simple machines belong to the **SISD** class (single instruction stream, single data stream), which means that all instructions are executed one at a time by a single processor, and each instruction is executed on a single piece of data. Actually, this is exactly our straightforward sequential programming model and does not involve any kind of parallelism.
- On a **MISD** type machine (multiple instruction stream, single data stream), different instructions can simultaneously operate on a single piece of data. Pipeline architectures belong to this type.
- In contrast, the **SIMD** class (single instruction stream, multiple data stream) is one of the two main architectures used for high performance computing. Representatives of this class are, basically, all vector computers. Vector operations are classical examples where such a machine is useful: a single instruction is performed for all components of a vector at once (e.g. the d multiplications required for computing $a \cdot x, x \in \mathbf{R}^d$ are computed simultaneously). Nowadays, many regular CPUs include SIMD instructions, at least on a small scale. Examples comprise SSE/AVX/AVX-2.
- The most important class in the supercomputing sector is that of the MIMD computers (multiple instruction stream, multiple data stream). This class can be considered to contain the "true" parallel computers, since here each processor has its own instruction and data streams. We can further discriminate within this category with respect to the way the memory is addressed:
 - computers with distributed memory (**MIMD distributed address space machines**) and
 - computers with a global memory (**MIMD global address space machines**).

For programming such machines, there are two important paradigms: either all nodes work on one (virtual) global memory or the nodes communicate via messages. The first paradigm fits perfectly to global address space architectures, whereas the latter one fits to distributed memory systems. Nevertheless, each programming paradigm might be applied on any architecture. **In this Practical Course, we shall concentrate on distributed memory machines programmed by message passing,** since on workstation or PC clusters, the distributed programming approach is most common, i.e. multiple parallel programs are run as processes or tasks on each processor. Data allocation occurs locally on each processor, and data exchange between individual processors is programmed explicitly with the help of respective instructions (send, receive, etc.). Programming of data communication is facilitated by communication libraries such as MPI (Message Passing Interface), which presents a




Figure 1: Domain decomposition methods with non-overlapping (**left**) and overlapping (**right**) partitions

hardware-independent interface. Respective libraries are offered by different vendors, some of them free of charge, and are available for Fortran, C, and C++. We explain the most important MPI commands throughout the next chapters and, especially, in the appendix.

3. Domain Decomposition as a Parallelization Strategy

A common approach to parallelize numerical algorithms for solving partial differential equations on MIMD architectures is to divide the computational domain Ω into subdomains $\Omega_1 \dots \Omega_N$ and to treat each subdomain by a separate process. The first so called *domain decomposition method* was developed in 1870 by H.A. Schwarz (a large class of domain decompositions is therefore referred to as Schwarz methods). Of course, in the 19th century, he did not invent this method to solve problems on a computer – instead, he was looking for means to find analytical solutions for PDEs.

In a strict sense, the term *domain decomposition* is used for a family of numerical methods that are based on solving a problem on each subdomain separately and combine these solutions to a global solution (in an iterative scheme, if necessary). Domain decomposition methods can be classified into overlapping and non-overlapping variants, depending on whether the individual subdomains overlap or not (see Figure 1). Moreover, we distinguish between **additive and multiplicative domain decomposition methods**. In multiplicative methods, the solutions of the subdomains are computed one after another, and the results from the previous subdomains may be used to compute that of the next domain. In contrast, additive methods compute the solution for all subdomains at once, and combine these solutions to obtain an update of the global solution. This option sounds especially attractive for parallelization!

Thus, each subdomain is assigned to a process that approximately computes the unknowns belonging to this subdomain. Each individual process, therefore, no longer requires access to the whole data set, but rather to a part of it. All processes can work in parallel on their part of the global data. The resulting **advantage is two-fold**: execution time is reduced because several processors are

sharing the work, and there is considerably more memory space available for the global problem than on sequential machines, because all processors can contribute their local memory.

The computing time is reduced most efficiently, if the computing work is equally distributed among the processors. In our application, this is accomplished when the subdomains each processor has to compute are of equal size and contain the same number of unknowns.

To ensure convergence of the algorithm on the entire domain, there has to be some communication between the domains. An exchange of relevant data between processors becomes necessary at certain times. In the case of domain decomposition, this concerns primarily the values of unknowns at the boundaries of neighboring processes. Exchange of data between processors, however, is an expensive operation, which can be several orders of magnitude slower than just copying data within main memory. Thus, the amount of data that has to be exchanged should be kept as small as possible, and the number of communication steps should be reduced to a minimum since establishing a connection between processes is typically expensive. Not paying attention to communication costs can completely annihilate the desired efficiency gain. Thus, one should try to minimize the domain borders, since the number of unknowns corresponds to the size of this border.

In a broader sense, any approach that distributes the computational domain of a problem to different processors is likely to be called a *domain decomposition* approach, even if it is not using the respective numerical schemes. The approach used in this worksheet mainly sticks to the idea of distributing the domain, too, and does not really make use of the numerical domain decomposition approach.

Part II.

Navier-Stokes in Parallel

4. Parallelization of the Flow Code

Let's turn to the parallelization of the sequential program described in Worksheet 1. We start by dividing the base domain Ω for the flow calculation along the grid lines into i_{proc} parts in the x direction and j_{proc} parts in the y direction, so that we get $i_{proc} \cdot j_{proc}$ rectangular subdomains $\Omega_{ip,jp}$, $ip = 1, \dots, i_{proc}$, $jp = 1, \dots, j_{proc}$. This decomposition is represented in Figure 2 where the grid lines of the staggered grid are shown as dashed lines and the subdomain boundaries as solid lines. In each process, the pressure and velocity values in the interior of the associated subdomain are calculated. To save an additional communication step, velocities lying exactly on a subdomain boundary are computed by both processes involved.

Thus, a process treating a subdomain $[(il - 1) \delta x, ir \delta x] \times [(jb - 1) \delta y, jt \delta y]$ requires the values shown in Figure 2. The values lying on points marked with white symbols must be computed by the process

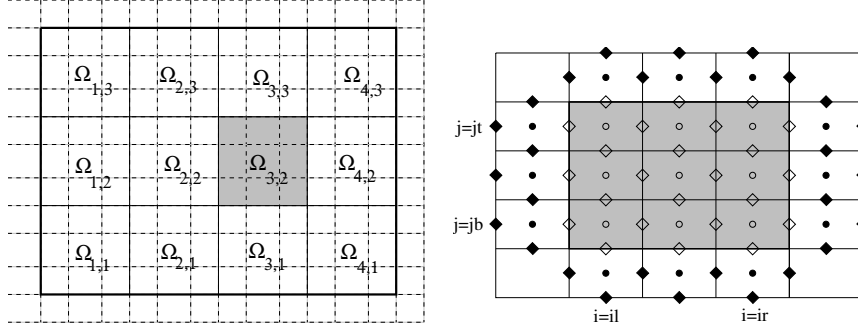


Figure 2: **Left:** Partitioning of the domain Ω ($\text{iproc}=4$, $\text{jproc}=3$). **Right:** Unknowns and boundary values for the subdomain $\Omega_{3,2}$

itself, whereas the values in the black points only serve as boundary data to determine the values in the interior and on the subdomain boundary. These boundary values are – unless they fall outside the total domain Ω – calculated by the neighboring processes. The boundary values must therefore be sent, in an appropriate form and with a proper frequency, by the neighboring processes. The following communication steps have to be implemented:

- After each SOR step in the pressure iteration, the pressure values located in the boundary strips must be exchanged as shown in Figure 3, such that the computation can proceed in the subdomains with the up-to-date boundary values. To avoid network overloading with too many simultaneous data bursts, the data exchange might be performed in four steps – to the left, to the right, up, and down. All relevant data should be sent to the respective neighbor as block of data.
- Following the data exchange of one SOR smoothing step, the termination criterion has to be evaluated: Therefore, each node computes its contribution to the global residual and, afterwards, sends the partial sum to a single process (master process). This process then decides whether to terminate the pressure iteration (tolerance ε achieved or the maximum number of iterations `itmax` exceeded) and broadcasts a message to all other processes saying whether to perform the next iteration step or to terminate the pressure iteration. Note that these steps can be simplified via an `MPI_Allreduce` call.
- At the end of the pressure iteration the current valid pressure values are contained in the boundary strips, so the velocity values on the subdomain boundaries can be updated in `calculate_uv` without any new communication. However, following `calculate_uv`, the updated velocity values must be exchanged as depicted in Figure 3 to enable the calculation of F and G in the next time step.
- Finally, the new time step `dt` must be determined. To find it, the maximum absolute values of U and V must be first determined in `calculate_dt` for each subdomain, and these local maxima should be sent to the master process. The latter determines the global maxima and disseminates them over all processes. After that, the CFL condition (see Worksheet 1, (13)) returns the same value of time step `dt` on each partial process. This procedure is, in principle, analogous to that of residual computation.

We may summarize all the steps in the following algorithm.




Figure 3: **Left:** Exchange of pressure values. **Right:** Exchange of velocity values

5. The Algorithm

```

Set  $t := 0, n := 0$ 
Assign initial values to  $u, v, p$ 
While  $t < t_{end}$ 
    Set boundary values for  $u$  and  $v$ 
    Compute  $F^{(n)}$  and  $G^{(n)}$  according to Worksheet 1
    Compute the right-hand side of the pressure equation (Worksheet 1)
    Set  $it := 0$ 
    While  $it < it_{max}$  and  $\|r^{it}\| > eps$ 
        Set the appropriate pressure boundary values (Worksheet 1)
        Perform an SOR cycle according to Worksheet 1
        Exchange the pressure values in the boundary strips
        Compute the partial residual sum and send it to the master process
        The master process computes the residual norm
            of the pressure equation  $\|r^{it}\|$  and broadcasts
            it to all the other processes (consider MPI_Allreduce)
         $it := it + 1$ 
    Compute  $u^{(n+1)}$  and  $v^{(n+1)}$  according to Worksheet 1
    Exchange the velocity values in the boundary strips
    Compute the maximum values of  $u^{(n+1)}$  and  $v^{(n+1)}$  for each
        process and send them to the master process (local max values)
    The master process computes the global maximum values of  $u^{(n+1)}$ 
        and  $v^{(n+1)}$  and broadcasts them to all other processes (consider MPI_Allreduce)
    Compute the common  $\delta t$  according to Worksheet 1
     $t := t + \delta t$ 
     $n := n + 1$ 

```

6. Implementation on Distributed Systems Using MPI

MPI stands for *Message Passing Interface* and defines a standard interface to incorporate many heterogeneous computers (workstations, parallel computers, vector machines) into a single virtual homogeneous parallel computer network. This standard has been established by the *MPI-Forum* (<https://www.mpi-forum.org>) and is also supported by a number of computer vendors providing their own libraries containing their native compilers and hardware extensions. Furthermore, there are also "free/open-source" libraries available (such as *MPICH* (<https://www.mpich.org/>) or *OpenMPI* (<https://www.open-mpi.org/>)) which we will use in this worksheet.

The parallel program development with MPI does not imply writing a different program for each process. Instead, one program is distributed among all the nodes, i.e. several instances of the program are running in the cluster (*SPMD* – single program multiple data, a special case of MIMD). To each program instance, a unique variable (process ID) is assigned and all the sender / receiver identification is done using these IDs.

In the parallelization of our flow code, we follow the principle of an additive, disjunct domain decomposition approach described in Section 4. Thus, for each process treating a subdomain with bounds $[(il - 1) \delta x, ir \delta x] \times [(jb - 1) \delta y, jt \delta y]$, the local arrays must have the following dimensions:

- P $[il-1, ir+1] \times [jb-1, jt+1]$
- U $[il-2, ir+1] \times [jb-1, jt+1]$
- V $[il-1, ir+1] \times [jb-2, jt+1]$ (*)
- F $[il-2, ir+1] \times [jb-1, jt+1]$
- G $[il-1, ir+1] \times [jb-2, jt+1]$
- RS $[il, ir] \times [jb, jt]$

Problems

1. Familiarize yourself with the concept of MPI-based parallel programming and MPI operations.
2. Get the files `parallel.h`, `parallel.c` from Moodle.
Furthermore, you need the files `init.h`, `init.c`, `boundary_val.h`, `boundary_val.c`, `uvp.h`, `uvp.c`, `visual.h`, `visual.c`, `helper.h`, `helper.c`, `main.c`.
3. Implement the following functions in `parallel.c` and `parallel.h`:
 - `void init_parallel (int iproc,int jproc,int imax,int jmax,
int *myrank,int *il,int *ir,int *jb,int *jt,
int *rank_l,int *rank_r,int *rank_b,int *rank_t,
int *omg_i,int *omg_j,int num_proc)`

The call of `MPI_Init(...)` must be present in the master program from the very beginning, so that in the routine to be written `n` processes should be taken into account.

At first, the boundaries `il,ir,jb,jt` and indices `omg_i, omg_j` for all subdomains must be computed in the master process and sent to the respectively assigned processes. For the computational domain of an arbitrary dimension and for an arbitrary number of processes (arranged in i- and j-directions), one must implement equipartitioning as uniform as possible.

Similarly, the neighborhood relationships should be determined, which means that the variables `rank_l,rank_r,rank_b,rank_t` should contain the process ID of the respective neighbors. The variables should be distributed for each subdomain, with the following MPI feature to be applied: if the subdomain lies, for example, at the physical boundary, then the process ID of a non-existent neighboring process `rank_l` to the left can be designated by the constant `MPI_PROC_NULL`. Should the process identified with `MPI_PROC_NULL` attempt later to send/receive data, then the corresponding MPI routine would return NOP (no operation). This feature spares some extraneous case differentiation, which makes the code more compact, transparent and efficient.

- `void pressure_comm(double **P,int il,int ir,int jb,int jt
int rank_l,int rank_r,int rank_b,int rank_t,
double *bufSend,double *bufRecv, MPI_Status *status, int chunk)`

This procedure exchanges pressure values between processes that treat adjacent subdomains. To evade a deadlock situation – for example, a situation where all processes are waiting to send or receive data, but to or from the wrong process –, one must stick to the following (or a similar) fixed order:

send to the left	— receive from the right,
send to the right	— receive from the left,
send to the top	— receive from the bottom,
send to the bottom	— receive from the top.

Instead of the trivial solution consisting in the `MPI_Send(...)/MPI_Recv(...)` combination, you should, whenever possible, use more efficient variants with `MPI_Sendrecv(...)`. Here, the vectors `bufSend, bufRecv` can serve as sending and receiving buffers. Be careful where/when to allocate the memory for your buffers `bufSend` and `bufRecv`, because these operations take a relatively long time.

This function is called within `sor` before actually computing the residual.

- `void uv_comm (double**U,double**V,int il,int ir, int jb,int jt,
int rank_l,int rank_r,int rank_b,int rank_t,
double *bufSend,double *bufRecv, MPI_Status *status, int chunk)`

This procedure is supposed to exchange the velocity values `U` and `V` between the processes treating adjacent subdomains. The data exchange should be performed in four steps as in `pressure_comm`, and `bufSend` as well as `bufRecv` have the same semantics as described

above. The values of U and V to be sent should be saved in a sufficiently large buffer `bufSend` or `bufRecv`. This function is called at the end of `calculate_uv`.

4. Perform the following modifications in the functions of Worksheet 1:

- All spatial operations (loops over i and j) must be carried out only in the respective subdomains, with the properties of staggered grids and boundary strips being taken into account.
- Since all informative `printf` outputs are executed by each process, the read-out feature is substantially limited. A remedy for this, especially for debugging, can be found in `parallel.c`, which provides some useful routines:
`Programm_Message(...)`, `Programm_Sync(...)`, `Programm_Stop(...)`.
 Another approach is to restrict the output to one process.
- In `visual.c`:

```
void output_uvp(double **U,double **V,double **P,int il,int ir,
               int jb,int jt,int omg_i,int omg_j,char *output_file)
```

Each individual process writes an output file of its own. Hereby, only elements in the domain of this process are to be plotted. Elements outside (in the strips, e.g.) shall not be written to the output file. Ensure, that the output filename for each process is unique. Paraview allows for the visualization of multiple input files. As the different domains are disjoint, our parallel visualisation requires for no special treatment: just load all the output files written by the visualization.

- In `read_parameter`:
 Enhance the function in such a way that the process partitioning arguments `iproc` and `jproc` can be additionally read out from the parameter file and passed to the main program.
- In `main`:
 While allocating the memory, the arrays must have dimensions given in (*). The main algorithm is, as described in Section 4, adjusted respectively. As the last instruction in the main program (before `return 0`), insert `Programm_Stop(...)` from the file `parallel.h`, which results in all processes to be synchronized and the `MPI_Finalize()` command to be called.
- In `boundary_values` and `spec_boundary_val`:
 The corresponding boundary values must be set when the subdomain boundary coincides with that of the whole domain.
- In `calculate_fg` and `sor`:
 The formulas for boundary values must be applied only if the subdomain boundaries coincide with those of the entire domain. Moreover, the calculation of the residual in SOR must be modified as described in Section 4; consider the MPI commands `MPI_Reduce` and `MPI_Allreduce` for this purpose.

- In `calculate_dt`:
For the adaptive time step control applied here, current values of maximal velocities over the entire computational domain are needed. As described in Section 4, a further communication step, similar to the one required for the residual calculation, has to be introduced.
- Test the parallel algorithm taking as an example the driven cavity with the following parameters:

```
imax = 300   jmax = 300   xlength = 10   ylength = 10
dt = 0.01    t_end = 1.0   tau = 0.5     dt_value = 2.0
eps = 0.01   omg = 1.7    alpha = 0.5   itermax = 100
GX = 0.0     GY = 0.0     Re = 10
UI = 0.0     VI = 0.0     PI = 0.0
iproc = 2    jproc = 3
```

7. Measuring Performance

If the grid is sufficiently fine and the number of processes is high enough, impressive computing speed increases may be achieved with a parallel program as compared to a serial program. This performance increase is usually expressed using the terms *speedup* and *efficiency*. These quantities are defined as

- speedup $S(p) := T(1)/T(p)$,
- parallel efficiency $E(p) := T(1)/(p \cdot T(p)) = S(p)/p$,

where p is the number of processes in use and $T(p)$ is the execution time of the parallel computation on p processes.

When measuring the runtimes on a busy network, though, neither the individual processors nor the computer network provide their resources exclusively to a single computing task. Instead, other processes compete for the available computing power and network bandwidth. Thus, for the set of parameters given above, i.e. a rather small problem, **it's also possible to gain only a small speed-up or even a slow-down, because the computational load is low in contrast to the time taken up by communication.** The advantages of parallelization really start to show in large computational problems. You should strive to achieve at least a noticeable speedup that improves when you increase the problem size on each partition. Otherwise, the purpose is primarily to get acquainted to the concepts and principal approaches to parallel systems and libraries as well as to illustrate these by example problems.

Part III.

Appendix

A. Parallelization using MPI

In order to facilitate studying MPI, a short program is presented in this appendix demonstrating the transition from code to running a program. Here the program `communication.c` is given, which writes a process ID (`rank`) for each process into `stdout`, a singled-out process (e.g. with ID 0) sends the value `PI` to all others, and finally each process notifies the neighboring processes (`myrank-1`, `myrank+1`) of its ID. Here we also show how the both boundary cases can be handled with the aid of a special MPI constant `MPI_PROC_NULL`.

```
#include <stdlib.h>
#include <stdio.h>
#include <mpi.h>

const double PI = 3.14;                                /* the value to be broadcasted */

int main(int argc, char* argv[]) {

    int i, myrank, nproc;
    int NeighborID_li, NeighborID_re;
    int recv_NeighborID_li, recv_NeighborID_re;
    double pi;
    MPI_Status status;

    /* initialisation */

    MPI_Init( &argc, &argv );                          /* execute n processes      */
    MPI_Comm_size( MPI_COMM_WORLD, &nproc );            /* asking for the number of processes */
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );           /* asking for the local process id */

    /* determine the neighbors */

    if ( 0 < myrank )
        NeighborID_li = myrank - 1;
    else
        NeighborID_li = MPI_PROC_NULL;                  /* no neighbor */

    if ( (nproc - 1) > myrank )
        NeighborID_re = myrank + 1;
    else
        NeighborID_re = MPI_PROC_NULL;                  /* no neighbor */
}
```

```

/* broadcast of pi from process 0 to all */

if ( 0 == myrank )
    pi = PI;
else
    pi = 0.0;

printf("Proc %2d : Before the broadcast pi = %e\n", myrank, pi);

MPI_Bcast( &pi, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD );

printf("Proc %2d : After the broadcast pi = %e\n", myrank, pi);

/* send the ID to left/right neighbors (conventional)
   think over how the following send/receive procedure can be implemented with
   the help of a single MPI_Sendrecv command */

/* ID of a non-existent neighbor has the value MPI_PROC_NULL */

MPI_Send( &myrank, 1, MPI_INT, NeighborID_li, 1, MPI_COMM_WORLD );
MPI_Recv( &recv_NeighborID_re, 1, MPI_INT, NeighborID_re, 1,
          MPI_COMM_WORLD, &status );

if ( MPI_PROC_NULL != NeighborID_re )
    printf("Proc %2d : ID right neighbor = %2d\n", myrank, recv_NeighborID_re);
else
    printf("Proc %2d : ID right neighbor = MPI_PROC_NULL\n", myrank);

MPI_Send( &myrank, 1, MPI_INT, NeighborID_re, 2, MPI_COMM_WORLD );
MPI_Recv( &recv_NeighborID_li, 1, MPI_INT, NeighborID_li, 2,
          MPI_COMM_WORLD, &status );

if ( MPI_PROC_NULL != NeighborID_li )
    printf("Proc %2d : ID left neighbor = %2d\n", myrank, recv_NeighborID_li);
else
    printf("Proc %2d : ID left neighbor = MPI_PROC_NULL\n", myrank);

/* Epilog */
fflush(stdout);
fflush(stderr);
MPI_Barrier( MPI_COMM_WORLD );
MPI_Finalize();

return 0;
}

```

Compilation and linking are done with `mpicc`, which accepts the general C-Compiler/Linker options. The following call can be used:

```
mpicc -o Communication Communication.c
```

Each MPI program is executed with the help of the `mpirun` shell script, to which at least the desirable number of processes as well as the executable should be passed as parameters:

```
mpirun -np 4 Communication
```

Then four `Communication` processes are started, which is quite sufficient during the development and debugging phase.

You may want to use the following MPI routines, types, and constants in this course:

```
MPI_Init, MPI_Finalize, MPI_Comm_rank, MPI_Comm_size, MPI_Barrier,  
MPI_Send, MPI_Recv, MPI_Sendrecv, MPI_Reduce, MPI_Allreduce, MPI_Bcast,  
MPI_COMM_WORLD, MPI_INT, MPI_DOUBLE, MPI_MAX, MPI_SUM,  
MPI_PROC_NULL
```

All the methods appearing in this worksheet are defined in the MPI Standard:

<http://www.mpi-forum.org/docs/docs.html>.

On Debian-based systems (such as Ubuntu), you may easily install OpenMPI using the packages `libopenmpi-dev` (headers and libraries) and `openmpi-bin` (executables).